

Going Deeper into Spark Core



Xavier Morera

HELPING DEVELOPERS UNDERSTAND SEARCH & BIG DATA

@xmorera www.xaviermorera.com



Going Deeper into Spark Core



Anonymous Functions / Lambdas

Named
Functions
in Spark

```
def split_the_line(x):  
    return x.split(',')
```

```
badges.map(add_one)
```



Why Are Lambdas so Useful?

Anonymous
Functions
in Spark

~~badges.map(add_one)~~

badges.map(lambda x: x.split(','))



You will find yourself using
lambdas all the time with Spark

Believe me...



```
spark2-submit --packages com.databricks:spark-xml_2.11:0.4.1,  
                    com.databricks:spark-csv_2.11:1.5.0  
                    prepare_data_posts_simple_titles.py
```

Extract Titles from Posts.xml



Data preparation step



```
lines = sc.textFile('/user/cloudera/stackexchange/simple_titles_txt')
words = lines.flatMap(lambda line: line.split(' '))
word_for_count = words.map(lambda x: (x,1))
word_for_count.reduceByKey(lambda x,y: x + y).collect()
```

A Closer Look at Map, FlatMap, Filter, Sort, ...

map() is one of the most commonly used transformations

Followed by **flatMap()**, **filter()** and **sort()**

And later on aggregations



```
word_for_count = words.map(lambda x: (x,1))  
word_for_count.take(1)  
words.map(lambda x: x.lower())  
words.map(lambda x: x.upper())
```


Map

Apply function to each element


map(), mapPartitions(), mapValues(), mapPartitionsWithIndex() ...

RDD of length N transformed to RDD of length N






Parent RDD



lambda x: x+1

Child RDD



```
word_for_count = words.map(lambda x: (x,1))  
word_for_count.take(1)  
words.map(lambda x: x.lower())  
words.map(lambda x: x.upper())
```

Map

Apply function to each element

map(), mapPartitions(), mapValues(), mapPartitionsWithIndex() ...

Each element in parent RDD mapped to one element in the child RDD



```
words = lines.flatMap(lambda line: line.split(' '))
```

FlatMap

Apply function to each element and returns list of elements

Returns 0, 1 or more elements, "flattens" the results with map



Parent RDD

How can I use DataFrames in 2.0
What is an RDD and Schema RDD
How do I group by a field
Can I use Hive from HUE

Child RDD



Parent RDD



Child RDD

How, can, I, use, DataFrames,
in, 2.0, What, is, an, RDD, and,
Schema, RDD, How, do, I,
group, by, a, field, Can, I, use,
Hive, from, HUE



```
def starts_h(word):  
    return word[0].lower().startswith('h')  
  
word_for_count.filter(starts_h).collect()
```

Filter

Apply a function to each element of the RDD

If the function returns false, element is not included in new RDD




Filter




Filter

Parent RDD



Child RDD



```
word_count = word_for_count.reduceByKey(lambda x,y: x + y)
word_count.take(10)

word_count.map(lambda
(x,y):(y,x)).sortByKey(False).map(lambda
(x,y):(y,x)).take(10)

word_count.sortBy(lambda (x,y): -y).take(10)
```

SortBy and SortByKey

Sort elements of an RDD

- By key on PairRDD with `sortByKey()`
- By a function using `sortBy()`



```
word_for_count.distinct().filter(starts_h).collect()
```

Many More Transformations

Plenty of transformations to go around

Some of them very powerful and/or very useful



Plenty of transformations to go around...



```
lines = sc.textFile('/user/cloudera/stackexchange/simple_titles.txt')
words = lines.flatMap(lambda line: line.split(' '))
word_for_count = words.map(lambda x: (x,1))
grouped_words = word_for_count.reduceByKey(lambda x,y: x + y)
grouped_words.collect()
```

Transformations

Start with method from **SparkContext** to load data

Transformations perform a computation

And create new RDDs



map(*f*, *preservesPartitioning=False*)

Return a new RDD by applying a function to each element of this RDD.

```
>>> rdd = sc.parallelize(["b", "a", "c"])
>>> sorted(rdd.map(lambda x: (x, 1)).collect())
[('a', 1), ('b', 1), ('c', 1)]
```

mapPartitions(*f*, *preservesPartitioning=False*)

Return a new RDD by applying a function to each partition of this RDD.

```
>>> rdd = sc.parallelize([1, 2, 3, 4], 2)
>>> def f(iterator): yield sum(iterator)
>>> rdd.mapPartitions(f).collect()
[3, 7]
```

mapPartitionsWithIndex(*f*, *preservesPartitioning=False*)

Return a new RDD by applying a function to each partition of this RDD, while tracking the index of the original partition.

```
>>> rdd = sc.parallelize([1, 2, 3, 4], 4)
>>> def f(splitIndex, iterator): yield splitIndex
```

Transformations

groupBy cartesian
flatMap intersection
map filter sample repartition union subtract
keyBy sortBy
coalesce zipWithIndex
zip mapPartitions
distinct



Transformations

PairRDDs

combineByKey sampleByKey

reduceByKey

join

leftOuterJoin

fullOuterJoin

sortByKey

cogroup

subtractByKey

groupByKey

rightOuterJoin

aggregateByKey

flatMapValues

foldByKey

reduceByKeyLocally

partitionBy




```
lines = sc.textFile('/user/cloudera/stackexchange/simple_titles.txt')
words = lines.flatMap(lambda line: line.split(' '))
word_for_count = words.map(lambda x: (x,1))
grouped_words = word_for_count.reduceByKey(lambda x,y: x + y)
grouped_words.collect()
```

Previously on Transformations

Transformations are what "changes" your data

Remember: Spark is lazy

No computation done when you specify transformation

A diagram consisting of a dotted circle with the word "words" inside it. To the right of the circle is the text ".collect()".



```
lines = sc.textFile('/user/cloudera/stackexchange/simple_titles_txt')
words = lines.flatMap(lambda line: line.split(' '))
word_for_count = words.map(lambda x: (x,1))
grouped_words = word_for_count.reduceByKey(lambda x,y: x + y)
grouped_words.collect()
```

Actions

Action triggers computation

i.e. can return data to the driver or save an RDD to storage

Operations that produce non RDD values



`collect()`

Return a list that contains all of the elements in this RDD.

Note: This method should only be used if the resulting array is expected to be small, as all the data is loaded into the driver's memory.

`collectAsMap()`

Return the key-value pairs in this RDD to the master as a dictionary.

Note: this method should only be used if the resulting data is expected to be small, as all the data is loaded into the driver's memory.

```
>>> m = sc.parallelize([(1, 2), (3, 4)]).collectAsMap()
>>> m[1]
2
>>> m[3]
4
```

Actions

histogram
saveAsHadoopDataset
collectAsMap
saveAsNewAPIHadoopDataset
collect
count
max
top
variance
mean
aggregate
fold
forEachPartition
saveAsHadoopFile
sampleVariance
countApprox
sum
first
takeSample
takeOrdered
treeAggregate
countApproxDistinct
saveAsTextFile
min
stdev
take
saveAsSequenceFile
saveAsObjectFile
treeReduce
saveAsNewAPIHadoopFile



Actions
PairRDD

countApproxDistinctByKey

keys
values

countByKey

countByValueApprox


countByKeyApprox

sampleByKeyExact

countByValue



A Thing or Two on Partitions



Partition is just a 'bunch' of data

One of the foundations of parallelism


Faster to operate within partition

- Than shuffling data

Group data to minimize network traffic



How Does Spark Partition Data?



Data locality


- Partition per HDFS block

Resources

Configuration or parameters



How Does Spark Partition Data?



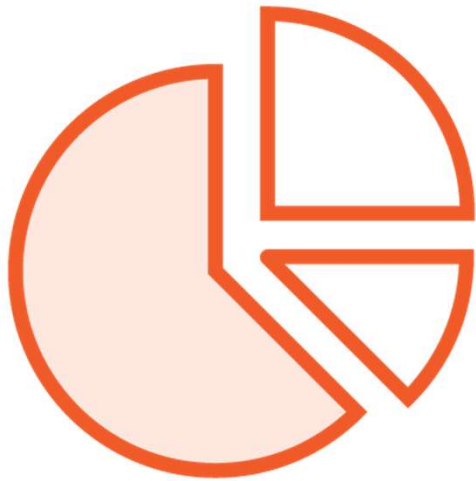
Partitioner

- Hash partitioner
- Range partitioner

Repartition



More or Less Partitions?



More partitions

- Less data per partition
- Smaller jobs
- More parallelism

Less partitions

- More data per partition
- Larger jobs



```
badges_for_part = badges_columns_rdd.map(lambda x:
(x[2],x)).repartition(50)

print badges_for_part.partitioner
def badge_partitioner(badge):
    return hash(badge)

badges_by_badge = badges_for_part.partitionBy(50,
badge_partitioner)
```

PartitionBy

Returns an RDD partitioned using a specific **partitioner**

Useful to get keyed data into same partition

Not yet a group operation



```
print badges_by_badge.partitioner
```

```
badges_for_part.saveAsTextFile('/user/cloudera/  
stackexchange/badges_nei_partitioner')
```

```
badges_by_badge.saveAsTextFile('/user/cloudera/  
stackexchange/badges_yei_partitioner')
```

PartitionBy

Create a function to be used for partitioning

Pass function as parameter to **partitionBy()**

Save with and without partitioner, and review results



```
for p in badges_by_badge.map(lambda (x,y):x)
    .glom().collect():
    print p
```

Glom

There is an action to coalesce all rows in a partition into an array

Useful for operations on all items within a partition

Let's print our keys per partition



Count Badges

```
def count_badges(iterator):  
    total = 0  
    for ite in iterator:  
        total += 1  
    yield total
```



```
counted_badges =  
badges_by_badge.mapPartitions(count_badges)  
counted_badges.collect()
```

MapPartitions

Apply a function to each partition

Done at a single pass

Returns after entire partition is processed



```
def next_value(value_list):  
    for I in value_list:  
        yield I  
  
test_yield = next_value([1, 2, 3])  
test_yield.next()
```

Yield

Returns a generator

Iterate with **next()**

Until **StopIteration**



```
posts_all.count()
```

Sampling Data

Selecting a representative part of the population

Faster, but you may lose accuracy

Also useful if you are resource constrained or very large dataset




```
posts_all.count()  
sample_posts=posts_all.sample(False, 0.1, 50)  
sample_posts.count()  
posts_all.countApprox(100, 0.95)
```

Sampling Data

Transformation to obtain a sample from your data with `sample()`

- withReplacement
- fraction
- seed



```
posts_all.count()  
sample_posts=posts_all.sample(False, 0.1, 50)  
sample_posts.count()  
posts_all.countApprox(100, 0.95)
```

Approximate Counts

Obtain an approximate count with **countApprox()**

Note: Experimental



```
posts_all.takeSample(False, 10, 50)  
len(posts_all.takeSample(False, 10, 50))
```

Take a Sample of Exact Size

Action available for exact count is called **takeSample()**



Set Operations

Questions

Answers



Set Operations

Posts Questions

('xavier', 1)


('troy', 2)

('xavier', 5)

Posts Answers

('xavier', 3)

('beth', 4)




```
questions=sc.parallelize([("xavier",1),("troy",2),  
("xavier",5)])  
  
answers=sc.parallelize([("xavier",3),("beth",4)])
```

Our Data

Create with parallelize

If you feel confident, go for the full dataset



```
questions.union(answers).collect()

questions.union(sc.parallelize(['irene', 'juli',
'luca'])).collect()

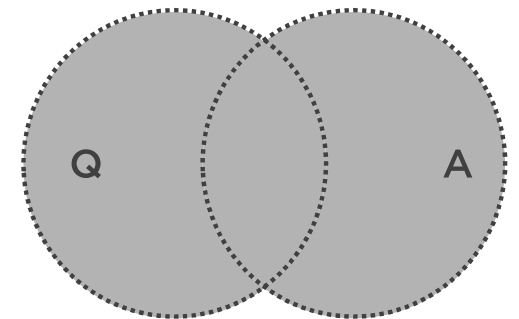
questions.union(sc.parallelize(range(10))).collect()
```


Union

RDD with all elements in both RDDs

Questions + answers

Can be different types





Union

All questions and answers

Elements remain the same





```
questions.join(answers).collect()  
questions.join(sc.parallelize(range(10))).collect()
```


Join

Elements with same keys in both, joined values

Hash join over the cluster, thus expensive

Unless known partitioner for narrow transformation





Join

People who have asked questions **AND** answered questions

Key is the person, value shows posts

Excludes those that do not contribute





```
questions.fullOuterJoin(answers).collect()
```

fullOuterJoin

Like **join()**, but....

None where key does not appear in one RDD





('xavier', (1, 3))

('xavier', (5, 3))

('troy', (2, None))

('beth', (None, 4))

fullOuterJoin

All questions and answers, joined by key

- None when user does not appear in one set





```
questions.leftOuterJoin(answers).collect()
```

leftOuterJoin

Join using keys from left set

None when key not found on right set





('beth', 4)

('xavier', (1, 3))

('xavier', (5, 3))

('troy', (2, None))

leftOuterJoin

Join using keys from left set

None when key not found on right set




```
questions.rightOuterJoin(answers).collect()
```

rightOuterJoin


Opposite of a leftOuterJoin

Join using keys from the right set

None where keys not available in left set



('troy', 2)



('xavier', (1, 3))

('xavier', (5, 3))

('beth', (None, 4))

rightOuterJoin

Opposite of a leftOuterJoin

Join using keys from the right set

None where keys not available in left set



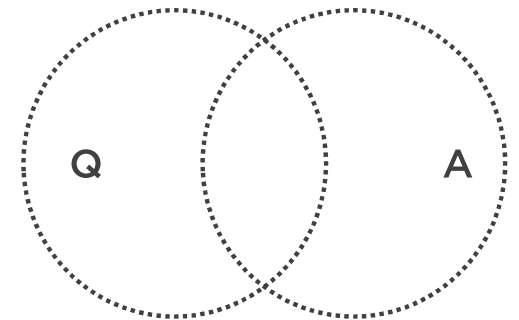

```
questions.leftOuterJoin(answers)  
answers.rightOuterJoin(questions)
```


leftOuterJoin and rightOuterJoin

```
questions.leftOuterJoin(answers)
```

Equivalent to

```
answers.rightOuterJoin(questions)
```





('beth', 4)

leftOuterJoin


('xavier', (1, 3))

('xavier', (5, 3))

('troy', (2, None))

rightOuterJoin

('beth', 4)



('xavier', (1, 3))

('xavier', (5, 3))


('troy', (2, None))

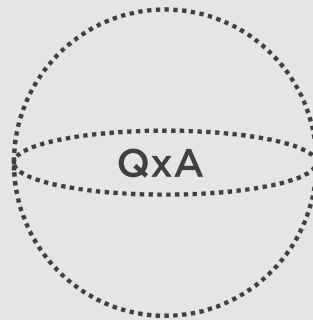
```
questions.cartesian(answers).collect()
```

Cartesian

Join of all elements in left set

With all elements in the right set





((('xavier', 1) ('xavier', 3)) (('xavier', 1) ('beth', 4)) (('troy', 2) ('xavier', 3))
 (('xavier', 5) ('xavier', 3)) (('troy', 2) ('beth', 4)) (('xavier', 5) ('beth', 4))

Cartesian

Join of all elements in left set

With all elements in the right set



Aggregation

Grouping elements together

Foundations of Big Data analytics



```
each_post_owner=posts_all.map(lambdax: x.split(",")[6])  
posts_owner_pair_rdd=each_post_owner.map(lambdax: (x,1))  
top_user_posts.map(lambda(x,y): (x,len(y))).take(1)
```

Prepare Some Data

Extract user from each post

PairRDD

- Key is user
- Value is 1



```
top_posters_gbk=posts_owner_pair_rdd.groupByKey()
```

GroupByKey

Values grouped by each key

Data sent over the network and collected on reduce workers

Can cause problems on larger datasets



```
top_user_posts = top_posters_gbk.filter(lambda (x,y):  
x == "51")
```

GroupByKey

Tuple of user id and list of 1's

Posts per user? → User id and number of posts

Use sortBy for top poster




```
from operator import add)
```

ReduceByKey

Perform an operation on all elements with same key

Specify a function

Reduce operation done within partition



```
top_posters_rbk=posts_owner_pair_rdd.reduceByKey(add)
top_posters_rbk.lookup('51')
top_posters_rbk.map(lambda(x,y):
(y,x)).sortByKey(False).map(lambda(x,y): (y,x)).take(10
```

ReduceByKey

Use **add**

Pass to **reduceByKey()**

Use **lookup()** to find top poster and confirm



```
top_posters_gbk.count()
```

```
top_posters_rbk.count()
```

groupByKey vs. reduceByKey

Do we get the same results?

Indeed we do



aggregateByKey

```
questions_asked=posts_all_entries.filter(lambda x:x[1]=="1")  
user_question_score=questions_asked.map(lambda x: (x[6],int(x[4])))  
for_keeping_count=(0,0)
```



```
aggregated_user_question=user_question_score.aggregateByKey(  
  
for_keeping_count, lambda tuple_sum_count, next_score:  
(tuple_sum_count[0]+next_score,  
  
tuple_sum_count[1]+1), lambda tuple_sum_count, tuple_next_partition_sum_c  
ount:(tuple_sum_count[0]+tuple_next_partition_sum_count[0],  
tuple_sum_count[1]+tuple_next_partition_sum_count[1]))
```

aggregateByKey

Like **reduceByKey()**

But takes an initial value

Specify functions for merging and combining



aggregateByKey

Combining

- Within partition

Merging

- Across partitions



```
aggregated_user_question.lookup( '51' )
```

aggregateByKey

Only questions, include score and user id

Define initial value, merging function, and combining function

Check with top poster



```
user_post = questions_asked.map(lambda x: (x[6],int(x[0])))
```

```
def to_list(postid):  
    return[postid]
```

```
def merge_posts(posta,postb):  
    posta.append(postb)  
    return posta
```

```
def combine_posts(posta, postb):  
    posta.extend(postb)  
    return posta
```




```
combined=user_post.combineByKey(to_list, merge_posts,  
combine_posts)  
combined.filter(lambda(x,y): x=='51').collect()  
combined.lookup('51')
```

CombineByKey

Specify an initial value can be a function that returns a new value

Provide merge and combine functions

Like **aggregateByKey()**, but more flexible



```
user_post.lookup( '51' )  
user_post.countByKey()[ '51' ]
```

CountByKey

Dictionary with keys and counts of occurrences

Like a **reduceByKey()** where we count based on key



reduceByKey & groupByKey

```
add_them = lambda x,y: x + y
add_in_list = lambda x: sum(list(x))
reduced = word_for_count.reduceByKey(add_them)
grouped =
word_for_count.groupByKey().mapValues(add_in_list)
```



```
reduced.take(1)
```

```
grouped.take(1)
```

```
grouped.count()
```

```
reduced.count()
```

reduceByKey & groupByKey

Both can be used for the same purpose

Aggregate by keys

Work very differently underneath



Comparing groupByKey vs. reduceByKey

groupByKey

| | | |
|-----------|--------------|--------------|
| (Spark,1) | (Spark,1) | (Cloudera,1) |
| (Spark,1) | (Spark,1) | (Cloudera,1) |
| (Spark,1) | (Spark,1) | |
| (Spark,1) | (Spark,1) | |
| (HUE,1) | (Cloudera,1) | |

reduceByKey

| | | |
|-----------|--------------|--------------|
| (Spark,1) | (Spark,1) | (Cloudera,1) |
| (Spark,1) | (Spark,1) | (Cloudera,1) |
| (Spark,1) | (Spark,1) | |
| (Spark,1) | (Spark,1) | |
| (HUE,1) | (Cloudera,1) | |



groupByKey

| | | |
|-----------|--------------|--------------|
| (Spark,1) | (Spark,1) | (Cloudera,1) |
| (Spark,1) | (Spark,1) | (Cloudera,1) |
| (Spark,1) | (Spark,1) | |
| (Spark,1) | (Spark,1) | |
| (HUE,1) | (Cloudera,1) | |

| | | |
|-----------|---------|--------------|
| (Spark,1) | (HUE,1) | (Cloudera,1) |
| (Spark,1) | | (Cloudera,1) |
| (Spark,1) | | (Cloudera,1) |
| (Spark,1) | | |
| (Spark,1) | | |
| (Spark,1) | | |
| (Spark,1) | | |

(Spark,8) (HUE,1) (Cloudera,3)

reduceByKey

| | | |
|----------------------|---------------------------|--------------|
| (Spark,4) (HUE,1) | (Spark,4) (Cloudera,1) | (Cloudera,2) |
|----------------------|---------------------------|--------------|

12 elements vs. 5 elements

(Spark,8) (HUE,1) (Cloudera,3)



Histogram



A diagram consisting of rectangles whose area is proportional to the frequency of a variable and whose width is equal to the class interval.



```
badges_reduced.map(lambda(x,y): y).histogram(7)
```

Grouping Data into Buckets with Histogram

Histograms are very powerful graphic tools

An image is worth a thousand words

Getting the data is usually the hardest part




```
badges_reduced.map(lambda(x,y): y)
                .histogram([0,1000,2000,3000,4000,5000,6000,7000])

badges_reduced.sortBy(lambda x:-x[1]).take(10)

badges_reduced.filter(lambda x: x[1]<1000).count()
```

Grouping Data into Buckets with Histogram

Specify number of intervals

- Returns array with intervals and array of counts within intervals

Explicitly state which intervals to use



Cache

Store data for future use, to improve response times

Persist to disk, memory or both



```
reduced.setName( 'Reduced RDD' )  
reduced.cache()
```

Cache & Persist

Spark may perform caching of intermediate results

- On expensive operations, to avoid recomputing when nodes fail



Details for Job 95

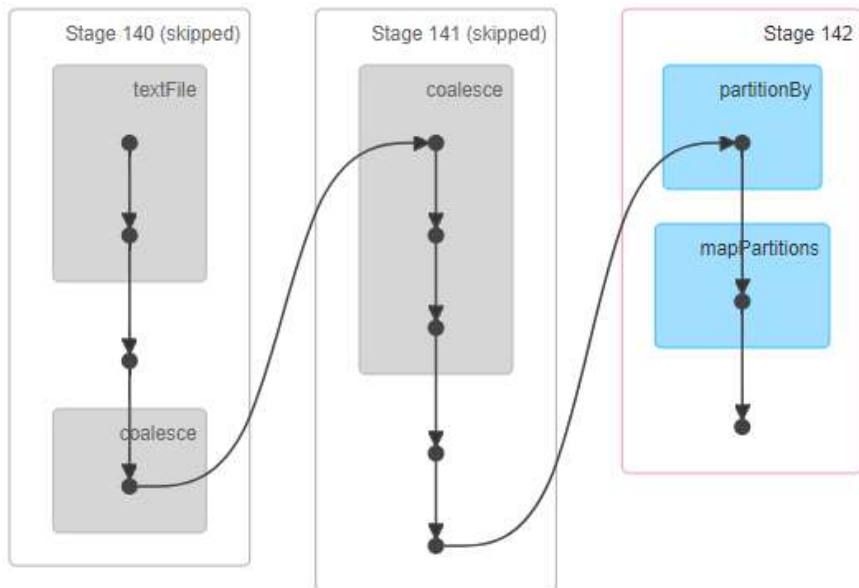
Status: SUCCEEDED

Completed Stages: 1

Skipped Stages: 2

▶ Event Timeline

▼ DAG Visualization



Completed Stages (1)

| Stage Id ▾ | Description | Submitted | Duration | Tasks: Succeeded/Total | Input | Output | Shuffle Read | Shuffle Write |
|------------|--|---------------------|----------|------------------------|-------|--------|--------------|---------------|
| 142 | runJob at PythonRDD.scala:446 +details | 2018/01/12 13:00:58 | 75 ms | 1/1 | | | 177.8 KB | |

```
badges_sorted.persist()
```

Cache & Persist

Spark may perform caching of intermediate results

- On expensive operations, to avoid recomputing when nodes fail

If the same job called twice, entire operation may be recomputed



```
grouped.setName( 'Grouped RDD' )  
grouped.persist(pyspark.storagelevel.StorageLevel.DISK_ONLY)
```

Cache & Persist


Call explicitly **cache()** and **persist()** when beneficial

- **cache()** is equivalent to **persist(MEMORY_ONLY)**

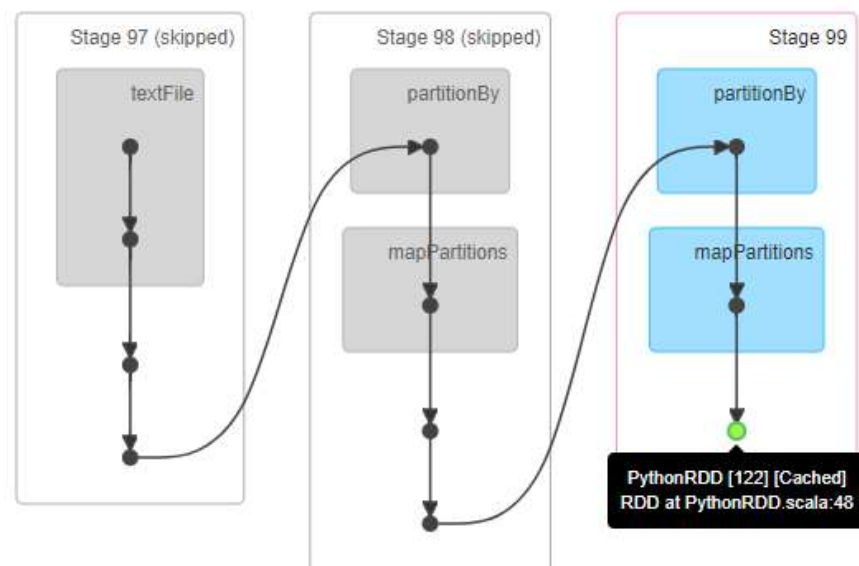

When RDD not needed anymore, call **unpersist()**




Storage Levels



Cache & Persist



Spark Processing



Distributed and parallel processing

Each executor has separate copies

- Variables and functions

No propagation data back to driver

- Except on certain necessary cases
- Accumulators & Broadcast Variables



Shared Variables

Accumulators

"Added"

Associate and commutative

Numeric accumulator

Other types possible

Counter is one common scenario

Accumulator may not be reliable

Case of failed task

Potential duplicate counts

Broadcast Variables

Read only variable

Immutable

Fits in memory

Distributed efficiently to the cluster

Do not modify after shipped

Good case is a lookup table



Accumulator

```
accumulator_badge=sc.accumulator(0)
```

```
accumulator_badge
```

```
def add_badge(item):
```

```
    accumulator_badge.add(1)
```

```
badges_by_badge.foreach(add_badge)
```



```
accumulator_badge.value
```

Accumulator

Create accumulator and check current value

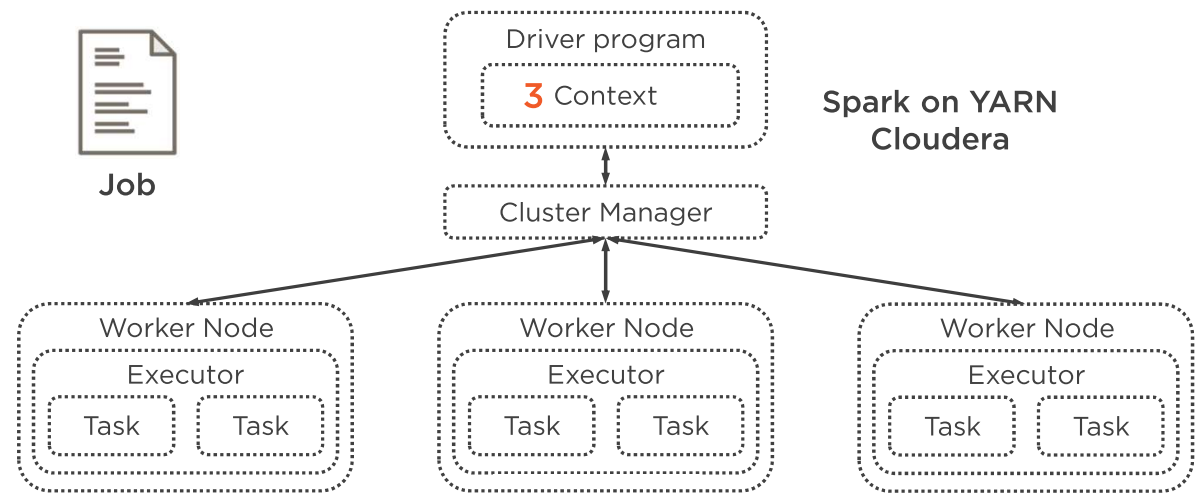
Increment accumulator function and run

Get value



Accumulators

Executors write to
accumulator in
Driver program



```
def get_name(user_column):  
    user_id = user_column[0]  
    user_name = user_column[3]  
    user_post_count = '0'  
    if user_id in broadcast_tp.value:  
        user_post_count = broadcast_tp.value[user_id]  
    return (user_id, user_name, user_post_count)
```

Broadcast Variable

Create a broadcast variable using the context

Access when necessary, i.e. lookup

Use `value`



```
tp = top_posters_rbk.collectAsMap()  
broadcast_tp = sc.broadcast(tp)  
user_info = users_columns.map(get_name)  
user_info.take(1)
```

Broadcast Variable

Create using **sc.broadcast()**

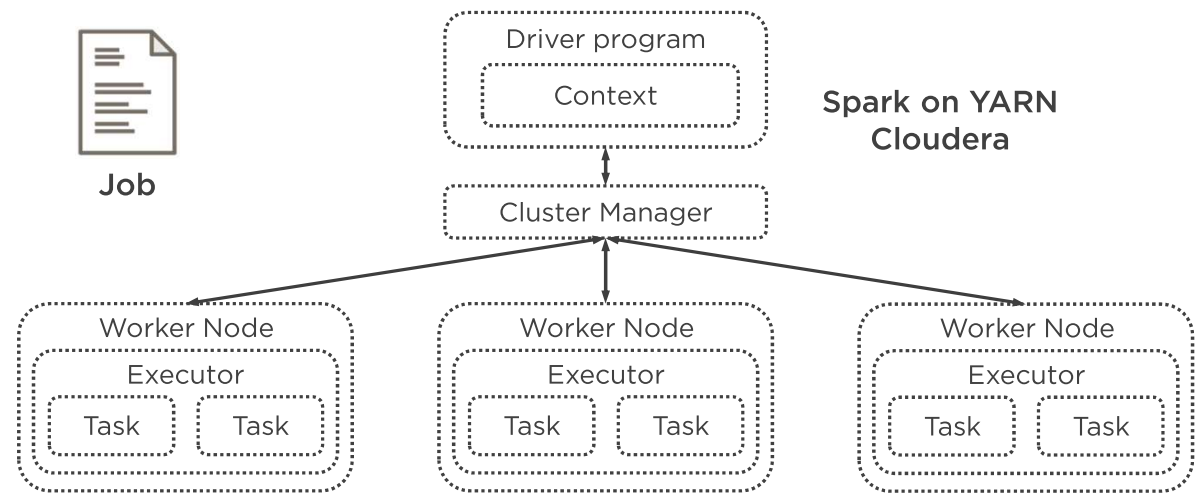
- Assign to a **variable**

Access using **variable.value**



Broadcast Variables

Executors read
from Broadcast
variable



Developing Self Contained PySpark Apps



Requires

- Create the SparkContext
- Dependencies
- Execute using spark2-submit



```
from pyspark import SparkContext  
sc = SparkContext("yarn", "Standalone App")
```

Creating the SparkContext

Corresponding **import**

Create **sc**



```
spark2-submit --py-files dependency.egg --jars ...
```

Dependencies

Use **py-files** for distributing files to cluster, i.e. zip file

Use also **jars** parameter

- Supports file, hdfs, http, ftp or local, but no directory expansion



```
spark2-submit <params-dependencies-conf> prepare_posts.py
```

Executing Application


Using **spark2-submit**

Pass any necessary configuration, dependencies and parameters

Code to be executed, submitted as a job



Disadvantages of RDDs



Don't take this the wrong way


RDDs are still used, even internally

Extremely powerful

Limitations on potential optimizations



Disadvantages of RDDs



Performance

Schema less

Steeper learning curve

"Everybody knows SQL"





Takeaway



Anonymous Functions

- Lambdas

Transformations vs. Actions

- Transformations return RDDs
- Actions trigger computation



Takeaway



Map, FlatMap, Filter, Sort, ...

Partitions

Sampling

Set operations

Aggregations



Takeaway



Histogram

Caching & Persisting

Shared variables

Self contained applications



Takeaway



Disadvantages of RDDs

| | | | |
|--|--|--|--|
| | | | |
| | | | |
| | | | |
| | | | |
| | | | |

